

Lecture 13:

Distance-vector Routing

CSE 123: Computer Networks
Alex C. Snoeren



HW 3 due now






Lecture 13 Overview

- Distance vector
 - ◆ Assume each router knows its own address and cost to reach each of its directly connected neighbors
- Bellman-Ford algorithm
 - ◆ Distributed route computation using only neighbor's info
- Mitigating loops
 - ◆ Split horizon and posion reverse



Bellman-Ford Algorithm

- Define distances at each node X
 - ♦ $d_x(y)$ = cost of least-cost *path* from X to Y
- Update distances based on neighbors
 - ♦ $d_x(y) = \min \{c(x,v) + d_v(y)\}$ over all neighbors V



$$d_u(z) = \min\{c(u,v) + d_v(z), \\ c(u,w) + d_w(z)\}$$



Distance Vector Algorithm

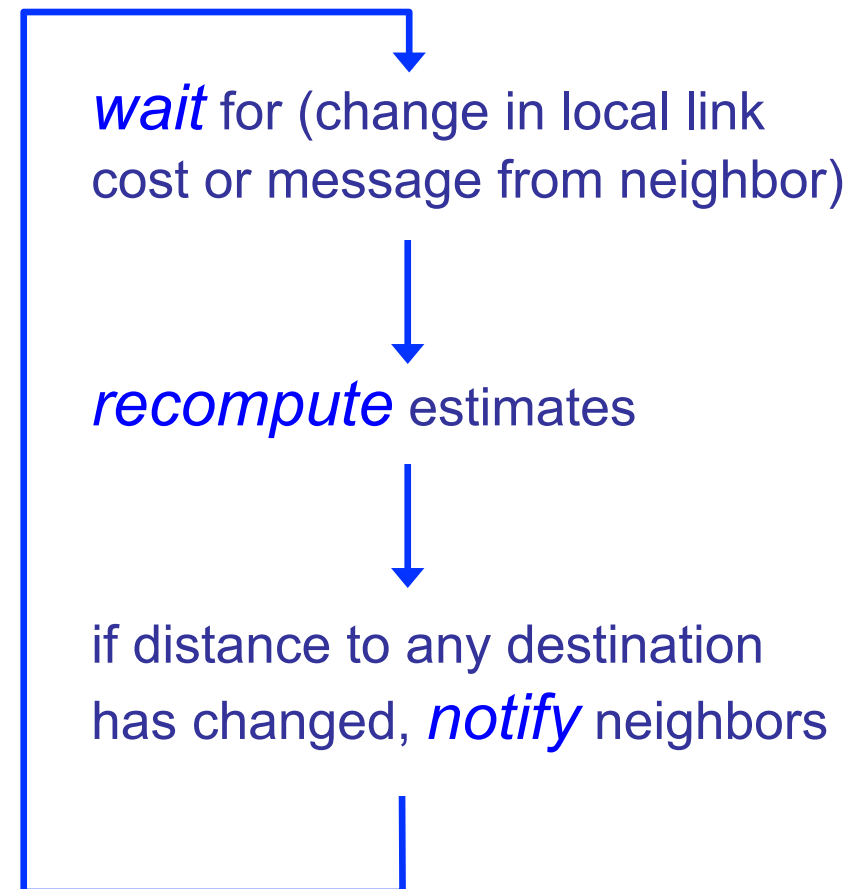
Iterative, asynchronous: each local iteration caused by:

- Local link cost change
- Distance vector update message from neighbor

Distributed:

- Each node notifies neighbors *only* when its DV changes
- Neighbors then notify their neighbors if necessary

Each node:






Step-by-Step

- $c(x, v)$ = cost for direct link from x to v
 - ♦ Node x maintains costs of direct links $c(x, v)$
- $D_x(y)$ = estimate of least cost from x to y
 - ♦ Node x maintains distance vector $\mathbf{D}_x = [D_x(y): y \in N]$
- Node x maintains its neighbors' distance vectors
 - ♦ For each neighbor v , x maintains $\mathbf{D}_v = [D_v(y): y \in N]$
- Each node v periodically sends \mathbf{D}_v to its neighbors
 - ♦ And neighbors update their own distance vectors
 - ♦ $D_x(y) \leftarrow \min_v \{c(x, v) + D_v(y)\}$ for each node $y \in N$




Example: Initial State



Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	∞	∞	1
B	7	0	1	∞	8
C	∞	1	0	2	∞
D	∞	∞	2	0	2
E	1	8	∞	2	0



D sends vector to *E*




Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	∞	∞	1
B	7	0	1	∞	8
C	∞	1	0	2	∞
D	∞	∞	2	0	2
E	1	8	4	2	0

D is 2 away, $2+2 < \infty$,
so best path to C is 4




B sends vector to A



Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	8	∞	1
B	7	0	1	∞	8
C	∞	1	0	2	∞
D	∞	∞	2	0	2
E	1	8	4	2	0



E sends vector to *A*




E is 1 away, $4+1 < 8$
so C is 5 away, $1+2 < \infty$
so D is 3 away

Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	5	3	1
B	7	0	1	∞	8
C	∞	1	0	2	∞
D	∞	∞	2	0	2
E	1	8	4	2	0

I'm 1 from A, 8 from B, 4 from C, 2 from D & 0 from E


...until Convergence



Info at node	Distance to Node				
	A	B	C	D	E
A	0	6	5	3	1
B	6	0	1	3	5
C	5	1	0	2	4
D	3	3	2	0	2
E	1	5	4	2	0



Node *B*'s distance vectors




Dest	Next hop		
	A	E	C
A	7	9	6
C	12	12	1
D	10	10	3
E	8	8	5



Handling Link Failure


- A marks distance to E as ∞ , and tells B
- E marks distance to A as ∞ , and tells B and D
- B and D recompute routes and tell C, E and E
- etc... until converge



Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	8	10	12
B	7	0	1	3	5
C	8	1	0	2	4
D	10	3	2	0	2
E	12	5	4	2	0



Counting to Infinity





Why so High?

- Updates don't contain enough information
- Can't totally order bad news above good news
- B accepts A 's path to C that is *implicitly* through B !
- Aside: this also causes delays in convergence even when it doesn't count to infinity



Mitigation Strategies



- **Hold downs**
 - ◆ As metric increases, delay propagating information
 - ◆ Limitation: Delays convergence
- **Loop avoidance**
 - ◆ Full path information in route advertisement
 - ◆ Explicit queries for loops (e.g. DUAL)
- **Split horizon**
 - ◆ Never advertise a destination through its next hop
 - » A doesn't advertise C to B
 - ◆ **Poison reverse:** Send negative information when advertising a destination through its next hop
 - » A advertises C to B with a metric of ∞
 - » Limitation: Only works for “loop”s of size 2



Poison Reverse Example


If Z routes through Y to get to X:

- Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)





Split Horizon Limitations



- A tells B & C that D is unreachable
- B computes new route through C
 - ♦ Tells C that D is unreachable (poison reverse)
 - ♦ Tells A it has path of cost 3 (split horizon doesn't apply)
- A computes new route through B
 - ♦ A tells C that D is now reachable
- Etc...

Routing Information Protocol



- DV protocol with hop count as metric
 - ◆ Infinity value is 16 hops; limits network size
 - ◆ Includes split horizon with poison reverse
- Routers send vectors every 30 seconds
 - ◆ With triggered updates for link failures
 - ◆ Time-out in 180 seconds to detect failures
- RIPv1 specified in RFC1058
 - ◆ www.ietf.org/rfc/rfc1058.txt
- RIPv2 (adds authentication etc.) in RFC1388
 - ◆ www.ietf.org/rfc/rfc1388.txt



Link-state vs. Distance-vector

Message complexity

- LS: with n nodes, E links, $O(nE)$ messages sent
- DV: exchange between neighbors only

Speed of Convergence

- LS: relatively fast
- DV: convergence time varies
 - ◆ May be routing loops
 - ◆ Count-to-infinity problem

Robustness: what happens if router malfunctions?

LS:

- Node can advertise incorrect *link* cost
- Each node computes only its *own* table

DV:

- Node can advertise incorrect *path* cost
- Each node's table used by others (error propagates)



Routing so far...

- Shortest-path routing
 - ◆ Metric-based, using link weights
 - ◆ Routers share a common view of path “goodness”
- As such, commonly used *inside* an organization
 - ◆ RIP and OSPF are mostly used as *intradomain* protocols
- But the Internet is a “network of networks”
 - ◆ How to stitch the many networks together?
 - ◆ When networks may not have common goals
 - ◆ ... and may not want to share information



For next time...

- Read Ch. 4.3.3-4 in P&D
- Keep moving on Project 2